

In [1]:

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt

from warnings import filterwarnings
filterwarnings(action='ignore')
```

In [2]:

```
thyroid_data = pd.read_csv("thyroid_data.csv")
thyroid_data
```

Out[2]:

	S.no	Age	Sex	On Thyroxine	Query on Thyroxine	On Antithyroid Medication	Sick	Pregnant	Thyroid Surgery	I131 Treatment
0	0	41	F	f	f	f	f	f	f	f
1	1	23	F	f	f	f	f	f	f	f
2	2	46	M	f	f	f	f	f	f	f
3	3	70	F	t	f	f	f	f	f	f
4	4	70	F	f	f	f	f	f	f	f
...
3216	2774	82	M	f	f	f	f	f	f	f
3217	2776	79	M	f	f	f	f	f	f	f
3218	2782	50	F	f	f	f	f	f	f	f
3219	2786	73	?	f	f	f	f	f	f	f
3220	2796	73	M	f	t	f	f	f	f	f

3221 rows × 28 columns

In [3]:

```
thyroid_data = thyroid_data.drop(['S.no'], axis = 1)
```

In [4]:

thyroid_data

Out[4]:

	Age	Sex	On Thyroxine	Query on Thyroxine	On Antithyroid Medication	Sick	Pregnant	Thyroid Surgery	I131 Treatment	Query Hypothyroid
0	41	F	f	f	f	f	f	f	f	f
1	23	F	f	f	f	f	f	f	f	f
2	46	M	f	f	f	f	f	f	f	f
3	70	F	t	f	f	f	f	f	f	f
4	70	F	f	f	f	f	f	f	f	f
...
3216	82	M	f	f	f	f	f	f	f	f
3217	79	M	f	f	f	f	f	f	f	f
3218	50	F	f	f	f	f	f	f	f	f
3219	73	?	f	f	f	f	f	f	f	f
3220	73	M	f	t	f	f	f	f	f	f

3221 rows × 27 columns

In [5]:

thyroid_data.shape

Out[5]:

(3221, 27)

In [6]:

```
## Columns
thyroid_data.columns
```

Out[6]:

```
Index(['Age', 'Sex', 'On Thyroxine', 'Query on Thyroxine',
      'On Antithyroid Medication', 'Sick', 'Pregnant', 'Thyroid Surgery',
      'I131 Treatment', 'Query Hypothyroid', 'Query Hyperthyroid', 'Lithium',
      'Goitre', 'Tumor', 'Hypopituitary', 'Psych', 'TSH Measured', 'TSH',
      'T3 Measured', 'T3', 'TT4 Measured', 'TT4', 'T4U Measured', 'T4U',
      'FTI Measured', 'FTI', 'Category'],
      dtype='object')
```

In [7]:

```
# A quick fix needed
thyroid_data.loc[thyroid_data['Age'] == '455', 'Age'] = '45'
```

In [8]:

```
## Let's drop some unnecessary columns
thyroid_data = thyroid_data.drop(['TSH Measured', 'T3 Measured', 'TT4 Measured', 'T4U Measured', 'FTI Measured', 'Category'])
```

In [9]:

```
#Checking for null values
thyroid_data.isna().sum()
```

Out[9]:

Age	0
Sex	0
On Thyroxine	0
Query on Thyroxine	0
On Antithyroid Medication	0
Sick	0
Pregnant	0
Thyroid Surgery	0
I131 Treatment	0
Query Hypothyroid	0
Query Hyperthyroid	0
Lithium	0
Goitre	0
Tumor	0
Hypopituitary	0
Psych	0
TSH	0
T3	0
TT4	0
T4U	0
FTI	0
Category	0

dtype: int64

In [10]:

```
thyroid_data.dtypes
```

Out[10]:

```
Age                object
Sex                object
On Thyroxine       object
Query on Thyroxine object
On Antithyroid Medication object
Sick               object
Pregnant           object
Thyroid Surgery    object
I131 Treatment     object
Query Hypothyroid  object
Query Hyperthyroid object
Lithium            object
Goitre             object
Tumor              object
Hypopituitary      object
Psych              object
TSH                object
T3                 object
TT4                object
T4U                object
FTI                object
Category           object
dtype: object
```

In [11]:

```
n = len(thyroid_data[thyroid_data['Category'] == 'hyperthyroid'])
print("No of hyperthyroid in Dataset:",n)
```

No of hyperthyroid in Dataset: 77

In [12]:

```
n1 = len(thyroid_data[thyroid_data['Category'] == 'hypothyroid'])
print("No of hypothyroid in Dataset:",n1)
```

No of hypothyroid in Dataset: 220

In [13]:

```
n2 = len(thyroid_data[thyroid_data['Category'] == 'sick'])
print("No of sick in Dataset:",n2)
```

No of sick in Dataset: 171


In [14]:

```
n3 = len(thyroid_data[thyroid_data['Category'] == 'negative'])
print("No of negative in Dataset:",n3)
```

No of negative in Dataset: 2753

In [15]:

```
fig = plt.figure()
ax = fig.add_axes([0,0,1,1])
ax.axis('equal')
l = ['hyperthyroid', 'hypothyroid', 'sick', 'negative']
s = [77, 220, 171, 2753]
ax.pie(s, labels = l, autopct='%1.2f%%')
plt.show()
```



In []:

In [16]:

```
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
from sklearn import svm
from sklearn import metrics
from sklearn.tree import DecisionTreeClassifier
```

In [17]:

```
train, test = train_test_split(thyroid_data, test_size = 0.20)
print(train.shape)
print(test.shape)
```

(2576, 22)

(645, 22)

In [18]:

```
thyroid_data.columns
```

Out[18]:

```
Index(['Age', 'Sex', 'On Thyroxine', 'Query on Thyroxine',  
      'On Antithyroid Medication', 'Sick', 'Pregnant', 'Thyroid Surgery',  
      'I131 Treatment', 'Query Hypothyroid', 'Query Hyperthyroid', 'Lithiu  
m',  
      'Goitre', 'Tumor', 'Hypopituitary', 'Psych', 'TSH', 'T3', 'TT4', 'T4  
U',  
      'FTI', 'Category'],  
      dtype='object')
```

In [19]:

```
thyroid_data.dtypes
```

Out[19]:

Age	object
Sex	object
On Thyroxine	object
Query on Thyroxine	object
On Antithyroid Medication	object
Sick	object
Pregnant	object
Thyroid Surgery	object
I131 Treatment	object
Query Hypothyroid	object
Query Hyperthyroid	object
Lithium	object
Goitre	object
Tumor	object
Hypopituitary	object
Psych	object
TSH	object
T3	object
TT4	object
T4U	object
FTI	object
Category	object
dtype:	object

In []:

In [20]:

```
def convert_category(dataframe, column):

    if column == 'Sex':
        conditionF = dataframe[column] == 'F' # For sex column
        conditionT = dataframe[column] == 'M' # For sex column
    else:
        conditionF = dataframe[column] == 'f'
        conditionT = dataframe[column] == 't'

    dataframe.loc[conditionF, column] = 0
    dataframe.loc[conditionT, column] = 1
```

In [21]:

```
# Binarize Category Columns
binary_cols = ['Age', 'Sex', 'On Thyroxine', 'Query on Thyroxine',
               'On Antithyroid Medication', 'Sick', 'Pregnant', 'Thyroid Surgery',
               'I131 Treatment', 'Query Hypothyroid', 'Query Hyperthyroid', 'Lithium',
               'Goitre', 'Tumor', 'Hypopituitary', 'Psych', 'TSH', 'T3', 'TT4', 'T4U',
               'FTI']

for col in binary_cols: convert_category(thyroid_data, col)
```

In [22]:

```
# Convert '?' to np.nan and convert numeric data to numeric dtype
for col in thyroid_data.columns:
    if col != 'Category':
        thyroid_data.loc[thyroid_data[col] == '?', col] = np.nan
        thyroid_data[col] = pd.to_numeric(thyroid_data[col])
```

In [23]:

```
from sklearn.impute import SimpleImputer

curr_columns = thyroid_data.columns.difference(['Category'])

imputer = SimpleImputer(missing_values=np.nan, strategy='median')
imputed_data = imputer.fit_transform(thyroid_data.drop('Category', axis=1))
imputed_data = pd.DataFrame(imputed_data, columns=curr_columns)
```

In [24]:

```
thyroid_data = pd.concat([
    imputed_data.reset_index(),
    thyroid_data['Category'].reset_index()],
    axis=1).drop('index', axis=1)
```

In [25]:

```
thyroid_data.dtypes
```

Out[25]:

```
Age                float64
FTI                float64
Goitre             float64
Hypopituitary      float64
I131 Treatment     float64
Lithium            float64
On Antithyroid Medication float64
On Thyroxine       float64
Pregnant           float64
Psych              float64
Query Hyperthyroid float64
Query Hypothyroid  float64
Query on Thyroxine float64
Sex                float64
Sick               float64
T3                float64
T4U               float64
TSH               float64
TT4               float64
Thyroid Surgery    float64
Tumor             float64
Category           object
dtype: object
```

In [26]:

```
thyroid_data.describe()
```

Out[26]:

	Age	FTI	Goitre	Hypopituitary	I131 Treatment	Lithium	Antith Medic
count	3221.000000	3221.000000	3221.000000	3221.000000	3221.000000	3221.000000	3221.00
mean	52.406085	0.306116	0.106489	0.014902	0.010866	0.043775	0.01
std	19.104151	0.460950	0.308510	0.121180	0.103689	0.204626	0.11
min	1.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.00
25%	37.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.00
50%	55.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.00
75%	68.000000	1.000000	0.000000	0.000000	0.000000	0.000000	0.00
max	94.000000	1.000000	1.000000	1.000000	1.000000	1.000000	1.00

8 rows × 21 columns

In [27]:

```

from sklearn.preprocessing import MinMaxScaler
from sklearn.model_selection import train_test_split

X = thyroid_data.drop('Category', axis=1)
y = thyroid_data['Category']

col_names = X.columns

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, stratify=y)

scaler = MinMaxScaler()
X_train = pd.DataFrame(scaler.fit_transform(X_train))
X_test = pd.DataFrame(scaler.transform(X_test))

```

In [28]:

```

#Using LogisticRegression
model = LogisticRegression()
model.fit(X_train, y_train)
prediction = model.predict(X_test)
print('Accuracy:', metrics.accuracy_score(prediction, y_test))

```

Accuracy: 0.8496124031007752

In [29]:

```

#Confusion matrix
from sklearn.metrics import confusion_matrix, classification_report
confusion_mat = confusion_matrix(y_test, prediction)
print("Confusion matrix: \n", confusion_mat)
print(classification_report(y_test, prediction))

```

Confusion matrix:

	precision	recall	f1-score	support
hyperthyroid	0.00	0.00	0.00	16
hypothyroid	0.50	0.05	0.08	44
negative	0.86	0.99	0.92	551
sick	0.00	0.00	0.00	34
accuracy			0.85	645
macro avg	0.34	0.26	0.25	645
weighted avg	0.77	0.85	0.79	645

In [30]:

```
#Using KNN Neighbors
from sklearn.neighbors import KNeighborsClassifier
model2 = KNeighborsClassifier(n_neighbors=5)
model2.fit(X_train,y_train)
y_pred2 = model2.predict(X_test)
```

In [31]:

```
from sklearn.metrics import accuracy_score
print("Accuracy Score:",accuracy_score(y_test,y_pred2))
```

Accuracy Score: 0.8232558139534883

In [32]:

```
#Using GaussianNB
from sklearn.naive_bayes import GaussianNB
model3 = GaussianNB()
model3.fit(X_train,y_train)
y_pred3 = model3.predict(X_test)
```

In [33]:

```
from sklearn.metrics import accuracy_score
print("Accuracy Score:",accuracy_score(y_test,y_pred3))
```

Accuracy Score: 0.10077519379844961

In [34]:

```
#Using Decision Tree
from sklearn.tree import DecisionTreeClassifier
model4 = DecisionTreeClassifier(criterion='entropy',random_state=7)
model4.fit(X_train,y_train)
y_pred4 = model4.predict(X_test)
```

In [35]:

```
from sklearn.metrics import accuracy_score
print("Accuracy Score:",accuracy_score(y_test,y_pred4))
```

Accuracy Score: 0.7162790697674418

In [36]:

```
# Fitting Naive Bayes Classification to the Training set with Linear kernel
from sklearn.naive_bayes import GaussianNB
nvclassifier = GaussianNB()
nvclassifier.fit(X_train, y_train)
```

Out[36]:

GaussianNB()

In [37]:

```
# Predicting the Test set results
y_pred = nvclassifier.predict(X_test)
print(y_pred)
```

```
['hypothyroid' 'hyperthyroid' 'hypothyroid' 'hyperthyroid' 'hypothyroid'
 'hyperthyroid' 'hyperthyroid' 'hyperthyroid' 'hypothyroid' 'hypothyroid'
 'hyperthyroid' 'hyperthyroid' 'hyperthyroid' 'hyperthyroid'
 'hyperthyroid' 'hyperthyroid' 'hyperthyroid' 'hypothyroid' 'hyperthyroi
d'
 'hyperthyroid' 'hyperthyroid' 'hypothyroid' 'hyperthyroid' 'hypothyroid'
 'hypothyroid' 'hyperthyroid' 'hyperthyroid' 'hyperthyroid' 'hypothyroid'
 'hyperthyroid' 'hypothyroid' 'hypothyroid' 'hyperthyroid' 'hyperthyroid'
 'hypothyroid' 'hyperthyroid' 'hyperthyroid' 'hypothyroid' 'hyperthyroid'
 'hyperthyroid' 'hyperthyroid' 'hyperthyroid' 'hyperthyroid'
 'hyperthyroid' 'hypothyroid' 'hyperthyroid' 'hyperthyroid' 'hypothyroid'
 'hyperthyroid' 'hyperthyroid' 'hyperthyroid' 'hyperthyroid'
 'hyperthyroid' 'hyperthyroid' 'hypothyroid' 'hypothyroid' 'hypothyroid'
 'hyperthyroid' 'hyperthyroid' 'hyperthyroid' 'hypothyroid' 'hypothyroid'
 'hyperthyroid' 'hypothyroid' 'hyperthyroid' 'hyperthyroid' 'hyperthyroi
d'
 'hyperthyroid' 'hyperthyroid' 'hypothyroid' 'hyperthyroid' 'hyperthyroi
d'
 'hypothyroid' 'hypothyroid' 'hyperthyroid' 'hypothyroid' 'hypothyroid'
 .....
```

In [38]:

```
#lets see the actual and predicted value side by side
y_compare = np.vstack((y_test,y_pred)).T
#actual value on the left side and predicted value on the right hand side
#printing the top 5 values
y_compare[:15,:]
```

Out[38]:

```
array([[ 'negative', 'hypothyroid'],
       [ 'negative', 'hyperthyroid'],
       [ 'negative', 'hypothyroid'],
       [ 'negative', 'hyperthyroid'],
       [ 'negative', 'hypothyroid'],
       [ 'negative', 'hyperthyroid'],
       [ 'hyperthyroid', 'hyperthyroid'],
       [ 'negative', 'hyperthyroid'],
       [ 'negative', 'hypothyroid'],
       [ 'hypothyroid', 'hypothyroid'],
       [ 'negative', 'hyperthyroid'],
       [ 'negative', 'hyperthyroid'],
       [ 'negative', 'hyperthyroid'],
       [ 'negative', 'hyperthyroid'],
       [ 'negative', 'hyperthyroid']], dtype=object)
```

In [39]:

```
# Making the Confusion Matrix
from sklearn.metrics import confusion_matrix
cm = confusion_matrix(y_test, y_pred)
print(cm)
```

```
[[ 16  0  0  0]
 [  0 40  1  3]
 [306 229  7  9]
 [ 15 17  0  2]]
```

In [40]:

```
#finding accuracy from the confusion matrix.
a = cm.shape
corrPred = 0
falsePred = 0

for row in range(a[0]):
    for c in range(a[1]):
        if row == c:
            corrPred +=cm[row,c]
        else:
            falsePred += cm[row,c]
print('Correct predictions: ', corrPred)
print('False predictions', falsePred)
print ('\n\nAccuracy of the Naive Bayes Clasification is: ', corrPred/(cm.sum()))
```

Correct predictions: 65
False predictions 580

Accuracy of the Naive Bayes Clasification is: 0.10077519379844961

In [41]:

```
#Using Support Vector
from sklearn.svm import SVC
model1 = SVC()
```

In [42]:

```
model1.fit(X_train,y_train)

pred_y = model1.predict(X_test)

from sklearn.metrics import accuracy_score
print("Acc=",accuracy_score(y_test,pred_y))
```

Acc= 0.8542635658914729

In [46]:

```
results = pd.DataFrame({
    'Model': ['Logistic Regression', 'KNN', 'Decision tree', 'GaussianNayeBayes', 'Support Vect
    'Score': [0.8496124031, 0.8232558, 0.7162790, 0.1007751, 0.85426356]})

result_df = results.sort_values(by='Score', ascending=False)
result_df = result_df.set_index('Score')
result_df.head(9)
```

Out[46]:

	Model
Score	
0.854264	Support Vector Machines
0.849612	Logistic Regression
0.823256	KNN
0.716279	Decision tree
0.100775	GaussianNayeBayes

In []:

In []:

In []:

In []: